

# Making Inheritance Work: C++ Issues

Steven Zeil

September 19, 2013

## Contents

<b>1</b>	<b>Base Class Function Members</b>	<b>2</b>
<b>2</b>	<b>Assignment and Subtyping</b>	<b>3</b>
<b>3</b>	<b>Virtual Destructors</b>	<b>4</b>
<b>4</b>	<b>Virtual Assignment</b>	<b>7</b>
<b>5</b>	<b>Virtual constructors</b>	<b>9</b>
5.1	Cloning . . . . .	11
<b>6</b>	<b>Downcasting</b>	<b>12</b>
6.1	RTTI . . . . .	13
<b>7</b>	<b>Single Dispatching &amp; VTables</b>	<b>14</b>
7.1	Single Dispatching . . . . .	15

## Recording

These slides accompany a recorded video: *Play Video*

.....

# 1 Base Class Function Members

## Base Class Function Members

Even if you override a function, the inherited bodies are still available to you.

```
class Person {
public:
    string name;
    long id;
    void print(ostream& out) {
        out << name << " " << id << endl;
    };

class Student: public Person {
public:
    string school;
    void print(ostream& out) { Person::print(out);
        out << school << endl;
    }
}
```

.....

## Base Class Constructors

This technique is often used in constructors so that subclasses will only need to initialize their own new data members:

```
class Person {
public:
    string name;
    long id;
    Person (string n, long i)
        : name(n), id(i)
    {}
};

class Student: public Person {
public:
```

```
string school;
Student (string name, long id, School s)
    : Person(name, id), school(s)
{}
}
```

.....

- This is a different use of initialization lists than we have seen before.
  - But is still consistent with the idea that the initialization list is actually a list of constructor calls.

## 2 Assignment and Subtyping

### Implementing Data Member Inheritance

Inheritance of data members is achieved by treating all new members as extensions of the base class:

```
class Person {
public:
    string name;
    long id;
};


class Student: public Person {
public:
    string school;
}
```

.....

### Extending Data Members

- When a compiler processes data member declarations, it assigns a byte offset to each one.

- In real life, these increase by however many bytes are required to store the previous data member. In this example, I'm going to pretend that each data member takes 4 bytes.



- Inherited members occur at the same byte offset as in the base class
- so code like `p->name` can translate the same whether `p` points to a `Person` or a `Student`.

## Making Inheritance Work: C++ Issues

- `p->name` is translated as “add 0 to the address in *p*”
- `p->id` is translated as “add 4 to the address in *p*”
- And that works for both Smith and Jones!

.....


### Assignment & Extension

In most OO languages, we can do

```
superObj = subObj;
```

but not

```
subObj = superObj;
```



- Assigning to a superclass object discards the extra data
  - Presumably, (Smith, 456782) is still a valid *Person*
  - Even if it loses the information about Smith being a student
    - \* So this at least can be said to make sense, even if it's not 100% safe.
- Assigning to a subclass object requires the system to invent data.
- If we assign Jones to a student object, what value should the system copy into the school?

.....

## 3 Virtual Destructors

### Virtual Destructors

As we've seen, subclasses can add new data members.

What happens if we add a pointer:

```
class Person {  
public:  
    string name;  
    long id;  
};  
  
class Student: public Person {
```

```
public:
    string school;
}

class GraduateStudent: public Student {
private:
    Transcript* undergradRecords;
public:
    ...
    GraduateStudent (const GraduateStudent& g);
    GraduateStudent& operator= (const GraduateStudent&);
    ~GraduateStudent();
};

GraduateStudent::GraduateStudent (const GraduateStudent& g)
    : name(g.name), id(g.id), school(g.school),
      undergradRecords(new Transcript(*(g.undergradRecords)))
{}

GraduateStudent& operator= (const GraduateStudent& g)
{
    if (this != &g)
    {
        Student::operator=(g);
        delete undergradRecords;
        undergradRecords = new Transcript(*(g.undergradRecords));
    }
    return *this;
}

GraduateStudent::~GraduateStudent()
{
    delete undergradRecords;
}
```

and we don't want to share?

.....

## Deleting Pointers and Inheritance

Consider the following two delete statements:

```
Person* g1 = new GraduateStudent(...);
GraduateStudent* g2 = new GraduateStudent(...);
:
delete g1; // compiler-generated ~Person() is called
delete g2; // ~GraduateStudent() is called
```

- Both calls are resolved by compile-time binding
  - Therefore the first delete leaks memory - undergraduateRecords is not cleaned up
- Fix would seem to be to force dynamic binding on the destructors

.....

### Making the Destructor Virtual

The trick is that this has to be done at the top of the inheritance hierarchy

```
class Person {
public:
    virtual ~Person() {}
    string name;
    long id;
};

class Student: public Person {
public:
    string school;
}

class GraduateStudent: public Student {
private:
    Transcript* undergradRecords;
public:
    :
    GraduateStudent (const GraduateStudent& g);
    GraduateStudent& operator= (const GraduateStudent&);
    ~GraduateStudent();
};
```

even though,

- at the time we wrote that class, there may have been no obvious need for a destructor
- this seems to violate the Rule of the Big 3
  - We'll look at the other two in just a moment

.....  
So you have to think ahead - if there's any chance of a non-shared pointer being added in a future subclass, make your destructor virtual.

## 4 Virtual Assignment

### Virtual Assignment

If subclasses can introduce new data members, should assignment be virtual so that we can guarantee proper copying of those extended data members?

### Virtual Assignment Example

```
void foo(Person& p1, const Person& p2)
{
    p1 = p2;
}

GraduateStudent g1, g2;
:
foo(g1, g2);
```

- If *p1* and *p2* "really" have *underGraduateRecord* fields, shouldn't we make sure those get copied properly during assignment?
  - Seems reasonable in this case.
  - But it means that assignment and copying will behave very differently, which is likely to catch programmers by surprise.

### What's the Problem with Virtual Assignment?

If you try it, the inherited members aren't what you might expect:

```
class Person {
public:
    virtual ~Person() {}
    virtual Person& operator= (const Person& p);
    string name;
    long id;
};

class Student: public Person {
public:
    string school;
    virtual Person& operator= (const Person& p); // inherited from Person
    // Student& operator= (const Student& s); // generated by compiler
}

class GraduateStudent: public Student {
private:
    Transcript* undergradRecords;
public:
    ...
    GraduateStudent (const GraduateStudent& g);
    virtual Person& operator= (const Person& p); // inherited from Person
    // Student& operator= (const Student& s); // inherited from Student
    GraduateStudent& operator= (const GraduateStudent&);
    ~GraduateStudent();
};
```

You actually wind up with multiple overloaded assignment operators in the subclasses.

.....

### What's the Problem with Virtual Assignment? (cont.)

- To make this work, you will need to implement both the virtual and the normal operators
- Implementing the virtual one is tricky because you might not get a GraduateStudent on the right:

```
void foo(Student& s1, const Student& s2)
{
```



```
s1 = s2;
}

Student s;
GraduateStudent g;
...
foo(g, s); // problem: s has no undergraduateRecords field
```

.....

### Recommendation

- There's no clear consensus in the C++ community about making assignment virtual.
- I recommend against it just because it's potentially confusing.
  - Try to avoid using assignment in situations where the "true" data type on the left is uncertain.

.....

## 5 Virtual constructors

### Virtual constructors

- Constructors can never be made virtual
- This can lead to problems when we need to create a copy.

.....

### Example: evaluating a cell reference

```
class CellReferenceNode: public Expression
{ // represents a reference to a cell
private:
    CellName value;

public:
    CellReferenceNode () {}
    <:::>

// Evaluate this expression
    virtual Value* evaluate(const Spreadsheet&) const;
    <:::>
```

*// Evaluate this expression*

```
Value* CellReferenceNode::evaluate(const Spreadsheet& s) const
{
    Cell* cell = s.getCell(value);
    Value* v = (Value*) cell->getValue();
    if (v == 0)
        return new ErrorValue();
    else
        return v;
}
```

- We would be better off returning a copy of the spreadsheet cell's value rather than the actual one.
  - Each Cell owns (does not share) its Value
  - Cell may therefore delete that Value
    - \* don't want to risk some other code doing so
- But how do we make a copy?

.....

### Not like this!

```
Value* theCopy = new Value(*v);
```

- How big is a Value?
- Would lose all data members in *v* required for its particular subtype of *Value*

.....

### Better, but Not the "OO Way"

```
Value* newCopy;
if (typeid(*v) == typeid(NumericValue)) {
    newCopy = new NumericValue (v->getNumericValue());
} else if (typeid(*v) == typeid(StringValue)) {
    newCopy = new StringValue (v->render(0));
} else if (typeid(*v) == typeid(ErrorValue)) {
    newCopy = new ErrorValue();
}
:
```

(We'll see how typeid works shortly.)

.....

## 5.1 Cloning

### Cloning

Solution is to use a simulated "virtual constructor", generally referred to as a `clone()` or `copy()` function.

```
Value* CellReferenceNode::evaluate(const Spreadsheet& s) const
{
    Cell* cell = s.getCell(value);
    Value* v = (Value*)cell->getValue();
    if (v == 0)
        return new ErrorValue();
    else
        return v->clone();
}
```

.....

### `clone()`

`clone()` must be supported by all values:

```
class Value {
public:
    :
    virtual Value* clone() const;
    :
};
```

.....

### Implementing `clone()`

Each subclass of `Value` implements `clone()` as a copy construction passed to `new`.

```
Value* NumericValue::clone() const
{
    return new NumericValue(*this);
}

Value* StringValue::clone() const
{
    return new StringValue(*this);
}
```

.....

## 6 Downcasting

Suppose that I want to be able to test any two Values to see if they are equal

- We'll define "equal" here as meaning that they are the same kind of value and would appear the same when rendered.

.....

### Example: Value::==

We want to explicitly require all subclasses of Value to provide this test:

```
class Value {
    :
    virtual bool operator== (const Value&) const;
};
```

The operator == compares two shapes. Its signature is: (const Value\*, const Value&) ⇒ bool

.....

### Inheriting ==

```
class NumericValue: public Value {
    :
class StringValue: public Value {
    :
    :
```

Both classes inherit the == operator. The signatures are

```
(const NumericValue*, const Value&) ⇒ bool
(const StringValue*, const Value&) ⇒ bool
```

.....

### Using the Inherited ==

```
NumericValue n1, n2;
StringValue s1, s2;
bool b = (n1 == n2)
        && (s1 == s2)
        && (n1 == s1);
```

The last clause suggests a problem.

- How do we compare values of different subtypes?
- Should we even allow it?

.....

### Implementing an asymmetric operator

We might implement `==` for `NumericValue` as:

```
bool NumericValue::operator==(
    (const Value& v)
{
    return d == v.d;
};
```

- But in a call like `(n1 == s1)`, `v.d` does not make sense.
  - In fact, this will get a compile error

.....

### Implementing an asymmetric operator (cont.)

The problem is that we can easily define

```
bool NumericValue::operator==(const NumericValue& v)
```

but

```
bool NumericValue::operator==(const Value& v)
```

seems impossible, as we cannot anticipate all the values that will ever be defined.

.....

## 6.1 RTTI

### Working around the `==` asymmetry

The C++ standard defines a mechanism for *RTTI* (Run Time Type Information).

```
bool NumericValue::operator==(const Value& v)
{
    if (typeid(v) == typeid(NumericValue)) {
        const NumericValue &nv =
            (const NumericValue&)v;
        return d == nv.d;
    } else
        return false;
};
```

- Note that `typeid()` can be applied both to objects and to types.
- But it can only be used with types/objects that have at least one virtual function.

.....

### RTTI: typeid and downcasting

RTTI also allows you to test to see if *v* is from a subclass of `NumericValue`

```
if (typeid(NumericValue).before(typeid(v))
```

or to perform safe *downcasting*:

```
NumericValue* np = dynamic_cast<NumericValue*>(&v);
if (np != 0)
{ // v really was a NumericValue or
  // subclass of NumericValue
  :
}
```

- The term “downcasting” refers to the fact that we are moving “down” in our inheritance hierarchy (assuming we draw the base class at the top).
  - Upcasting is always safe (and usually is done implicitly)
  - Downcasting can be dangerous if we don’t check to see if the object really is what we think it will be.

.....

### Downcasting Should Not Be a Crutch

Downcasting is often a tempting way to patch a poor initial choice of virtual “protocol” functions.

- 95% of the time, it’s a bad idea
  - often leads to subtle, hard to trace bugs

Oddly, though, downcasting is far more widely accepted in Java than in C++.

.....

## 7 Single Dispatching & VTables

### Equality Again

Earlier, we looked at the problem of comparing two spreadsheet `Value`s:

```
class Value {
:
virtual bool isEqual (const Value&) const;
};
```

```
class NumericValue: public Value {
    :
    virtual bool isEqual (const Value&) const;
};
```

We saw that problems are caused by `NumericValue::isEqual` getting a parameter of type `Value&` rather than `NumericValue&`.

.....

### Why is this so hard?

Why can't we select the best fit from among:

```
class NumericValue: public Value {
    :
    virtual bool isEqual (const NumericValue&) const;
    virtual bool isEqual (const StringValue&) const;
    virtual bool isEqual (const ErrorValue&) const;
};
```

The answer stems from how dynamic binding is implemented.

.....

## 7.1 Single Dispatching

### Single Dispatching

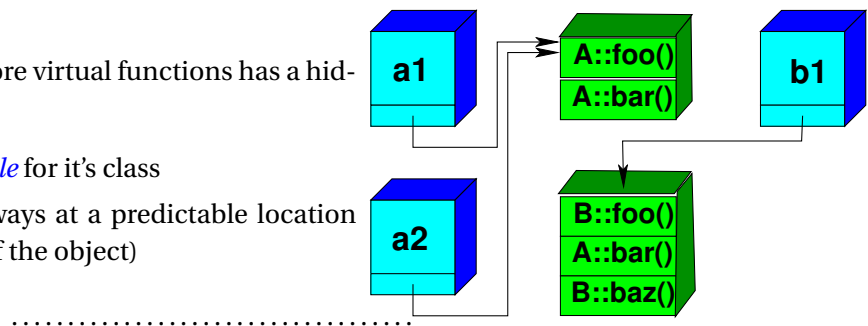
Almost all OO languages offer a *single dispatch* model of message passing:

- the dynamic binding is resolved according to the type of the single object to which the message was sent ("dispatched").
  - In C++, this is the object on the left in a call: `obj.foo(...)`
- There are times when this is inappropriate.
  - But it leads to a fast, simple implementation

.....

### VTables

- Each object with 1 or more virtual functions has a hidden data member.
  - a pointer to a *VTable* for its class
  - this member is always at a predictable location (e.g., start or end of the object)



### Compiling Virtual Function Declarations

- Each virtual function in a class is assigned a unique, consecutive index number.
- $(*VTable)[i]$  is the address of the class's method for the  $i$ 'th virtual function.

.....

### Example of VTable Use

```
class A {
public:
    A();
    virtual void foo();
    virtual void bar();
};

class B: public A {
public:
    B();
    virtual void foo();
    virtual void baz();
};
```

```
A* a = ???; // might point to an A or a B object
a->foo();
```

`foo()`, `bar()`, and `baz()` are assigned indices 0, 1, and 2, respectively.

.....

### Example: VTable Structure



## Making Inheritance Work: C++ Issues


- The call `a->foo()` is translated as

```
*(a->VTABLE[0])();
```

- The call `a->bar()` is translated as

```
*(a->VTABLE[1])();
```

Notice that this works regardless of whether *a* points to an *A* object or a *B* object.



- “works” in this case means “does dynamic binding”

.....

- Note that the call `a->baz()` would not compile, so we should not have to worry about going past the end of the shorter vtable.

## Implementing RTTI

- The address of the VTable is a unique identifier for each class known to the compiler
- This makes the vtable an ideal for implementing RTTI
  - and explains why RTTI is only available for classes with at least one virtual function

.....